

LAB Assignment No. 5

Topic: Artificial Neural Network Model

Question 1

Logic Gates with Neural Network. Implement a feed-forward neural network to learn the AND gate.

- Inputs: (0,0), (0,1), (1,0), (1,1)
- Output: 0, 0, 0, 1

Tasks:

1. Create dataset using NumPy or pandas.
2. Build a neural network with one hidden layer using TensorFlow/Keras or PyTorch.
3. Train it and show accuracy.
4. Compare model predictions with actual outputs

The screenshot shows a Jupyter Notebook interface with three code cells and their corresponding outputs.

- Cell 50:** Contains Python code to import numpy and pandas, and to define X and y arrays. The output shows the arrays have been created successfully.
- Cell 51:** Contains Python code to create a DataFrame from the X and y arrays. The output shows the DataFrame has been printed.
- Cell 52:** Contains Python code to print the DataFrame. The output shows the DataFrame with four rows and four columns: Input1, Input2, and Output.

	Input1	Input2	Output
0	0	0	0
1	0	1	0
2	1	0	0
3	1	1	1

```

model = Sequential()
model.add(Dense(4, input_dim=2, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

[53] ✓ 0.0s Python
... c:\Users\h\AppData\Local\Programs\Python\Python313\Lib\site-packages\keras\src\layers\core\dense.py:95: UserWarning: Do not pass an `input_shape` / `input_dim` argument to `Dense`'s constructor. Instead, pass an integer to `input_dim` or a tuple/list to `input_shape` (e.g., (batch_size, input_dim)). If you are passing a `InputSpec` object, pass it to the constructor instead. See: https://keras.io/api/layers/core_layers/dense/#dense-class
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)

model.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy']
)

[54] ✓ 0.0s Python

model.fit(X, y, epochs=500, verbose=0)

[55] ✓ 55.0s Python
... <keras.src.callbacks.history.History at 0x15e27566780>

loss, accuracy = model.evaluate(X, y, verbose=0)
print("Model Accuracy:", accuracy)

[56] ✓ 0.3s Python
... Model Accuracy: 1.0

```



```

+ labs.ipynb > comparison = pd.DataFrame({
+ Code + Markdown | ▶ Run All ⚡ Restart ⌛ Clear All Outputs | Jupyter Variables ⌂ Outline ...
  loss, accuracy = model.evaluate(X, y, verbose=0)
  print("Model Accuracy:", accuracy)

[56] ✓ 0.3s Python
... Model Accuracy: 1.0

predictions = model.predict(X)

[57] ✓ 0.1s Python
... 1/1 ━━━━━━ 0s 108ms/step

predicted_output = (predictions > 0.5).astype(int)

[58] ✓ 0.0s Python

comparison = pd.DataFrame({
    "Input1": X[:, 0],
    "Input2": X[:, 1],
    "Actual Output": y,
    "Predicted Output": predicted_output.flatten()
})

print(comparison)

[59] ✓ 0.0s Python
...   Input1  Input2  Actual Output  Predicted Output
  0       0       0             0             0
  1       0       1             0             0
  2       1       0             0             0
  3       1       1             1             1

```

Question 2

Create a dataset $y = x^2 + \text{noise}$ for x in range $[-3,3]$. Regression Task with Neural Network

Tasks:

1. Generate 100 samples.
2. Build a neural network to predict y from x.
3. Plot actual vs. predicted results.
4. Discuss how increasing hidden neurons changes results.

The screenshot shows a Jupyter Notebook interface with three code cells. The first cell (cell 87) contains code to generate 100 samples of data. The second cell (cell 88) contains code to define the network architecture, initialize weights and biases, and define the ReLU activation function and its derivative. The third cell (cell 89) is currently being edited.

```
import numpy as np
import matplotlib.pyplot as plt

[87] ✓ 0.0s Python

np.random.seed(0)

# 100 samples in range [-3,3]
x = np.linspace(-3, 3, 100).reshape(-1, 1)
noise = np.random.randn(100,1) * 0.5
y = x**2 + noise

[88] ✓ 0.0s Python

# Network architecture
input_neurons = 1
hidden_neurons = 8
output_neurons = 1

# Initialize weights and biases
W1 = np.random.randn(input_neurons, hidden_neurons)
b1 = np.zeros((1, hidden_neurons))

W2 = np.random.randn(hidden_neurons, output_neurons)
b2 = np.zeros((1, output_neurons))

# Activation function
def relu(z):
    return np.maximum(0, z)

def relu_derivative(z):
    return (z > 0).astype(float)

[89]
```

```

+ Code + Markdown | ▶ Run All ⌂ Restart ⌂ Clear All Outputs | Jupyter Variables | Outline ...
learning_rate = 0.01
epochs = 2000

for epoch in range(epochs):
    # Forward pass
    z1 = x @ W1 + b1
    a1 = relu(z1)
    z2 = a1 @ W2 + b2
    y_pred = z2 # Linear output for regression

    # Loss (MSE)
    loss = np.mean((y_pred - y)**2)

    # Backpropagation
    dloss = 2*(y_pred - y)/y.size
    dw2 = a1.T @ dloss
    db2 = np.sum(dloss, axis=0, keepdims=True)

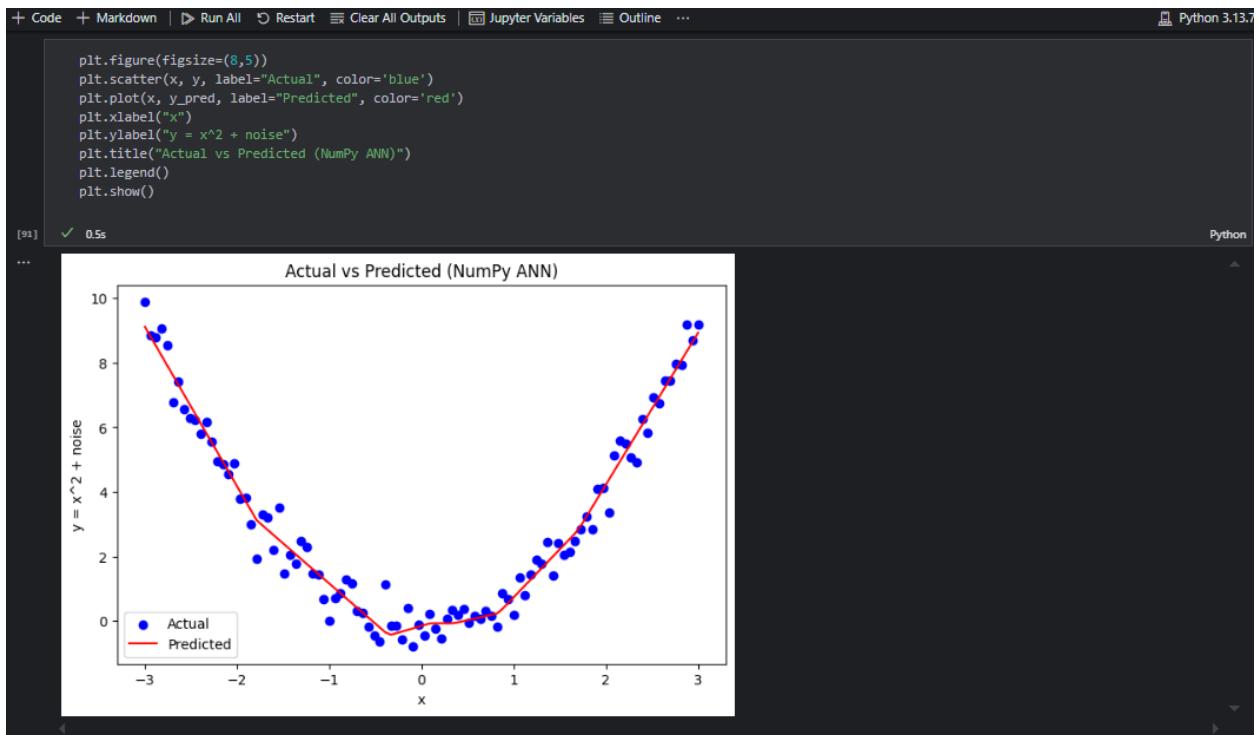
    da1 = dloss @ W2.T
    dz1 = da1 * relu_derivative(z1)
    dw1 = x.T @ dz1
    db1 = np.sum(dz1, axis=0, keepdims=True)

    # Update weights
    W2 -= learning_rate * dw2
    b2 -= learning_rate * db2
    W1 -= learning_rate * dw1
    b1 -= learning_rate * db1

    if (epoch+1) % 500 == 0:
        print(f"Epoch {epoch+1}, Loss: {loss:.4f}")

[90] ✓ 0.5s Python
...
Epoch 500, Loss: 0.3950
Epoch 1000, Loss: 0.2915
Epoch 1500, Loss: 0.2392
Epoch 2000, Loss: 0.2207

```



Question 3:

Use the XOR gate and train networks with different activation functions (sigmoid, tanh, ReLU).

- Compare accuracy, loss, and convergence speed.

- Plot and discuss results.

The screenshot shows a Jupyter Notebook interface with the following details:

- Header:** lab05.ipynb > plt.figure(figsize=(10,5))
- Toolbar:** Code | Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline | ...
- Version:** Python 3.13.7
- Cells:**
 - [110] 0.0s Python:

```
import numpy as np
import matplotlib.pyplot as plt
```
 - [111] 0.0s Python:

```
# Inputs
X = np.array([[0,0],[0,1],[1,0],[1,1]])
y = np.array([[0],[1],[1],[0]])
```
 - [112] 0.0s Python:

```
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(x):
    s = sigmoid(x)
    return s * (1 - s)

def tanh(x):
    return np.tanh(x)

def tanh_derivative(x):
    return 1 - np.tanh(x)**2

def relu(x):
    return np.maximum(0, x)

def relu_derivative(x):
    return (x > 0).astype(float)
```

```

def train_xor(activation, activation_derivative, hidden_neurons=4, lr=0.1, epochs=10000):
    # Initialize weights
    np.random.seed(0)
    W1 = np.random.randn(2, hidden_neurons)
    b1 = np.zeros((1, hidden_neurons))
    W2 = np.random.randn(hidden_neurons, 1)
    b2 = np.zeros((1,1))

    losses = []

    for epoch in range(epochs):
        # Forward pass
        z1 = X @ W1 + b1
        a1 = activation(z1)
        z2 = a1 @ W2 + b2
        y_pred = sigmoid(z2) # Output layer always sigmoid

        # Loss
        loss = np.mean((y_pred - y)**2)
        losses.append(loss)

        # Backpropagation
        dloss = 2*(y_pred - y)/y.size
        dz2 = dloss * sigmoid_derivative(z2)
        dW2 = a1.T @ dz2
        db2 = np.sum(dz2, axis=0, keepdims=True)

        da1 = dz2 @ W2.T
        dz1 = da1 * activation_derivative(z1)
        dW1 = X.T @ dz1
        db1 = np.sum(dz1, axis=0, keepdims=True)

        # Update weights
        W2 -= lr * dW2
        b2 -= lr * db2
        W1 -= lr * dW1
        b1 -= lr * db1

```

```

+ Code + Markdown | ▶ Run All ⏪ Restart ⏹ Clear All Outputs | Jupyter Variables ⏹ Outline ...
Python 3.13.7

    # Update weights
    W2 -= lr * dW2
    b2 -= lr * db2
    W1 -= lr * dW1
    b1 -= lr * db1

    # Final predictions
    final_pred = (y_pred > 0.5).astype(int)
    accuracy = np.mean(final_pred == y)
    return losses, accuracy, final_pred

[113] ✓ 0.0s Python

# Sigmoid
loss_sigmoid, acc_sigmoid, pred_sigmoid = train_xor(sigmoid, sigmoid_derivative)

# Tanh
loss_tanh, acc_tanh, pred_tanh = train_xor(tanh, tanh_derivative)

# ReLU
loss_relu, acc_relu, pred_relu = train_xor(relu, relu_derivative)

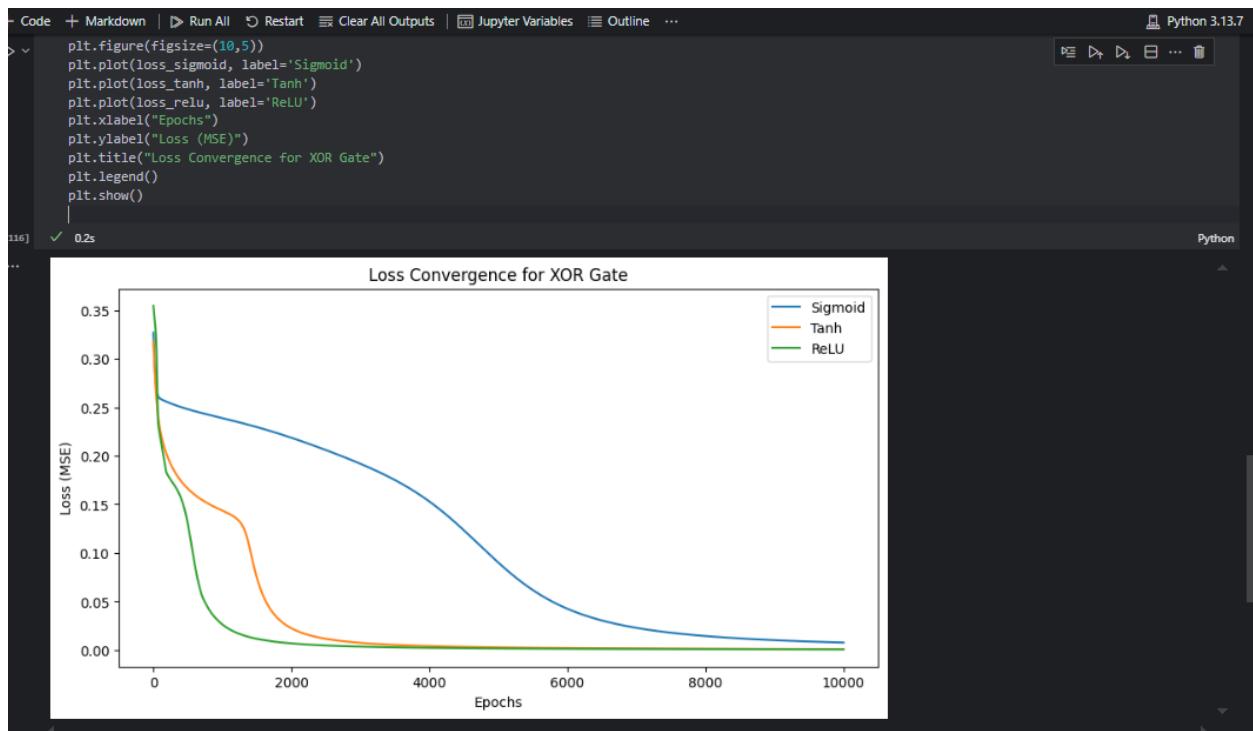
[114] ✓ 52s Python

▷ 
print("Accuracy with Sigmoid:", acc_sigmoid)
print("Accuracy with Tanh   :", acc_tanh)
print("Accuracy with ReLU   :", acc_relu)

[115] ✓ 0.0s Python

... Accuracy with Sigmoid: 1.0
Accuracy with Tanh   : 1.0
Accuracy with ReLU   : 1.0

```



Question 4

Binary Classification using Neural Network

Objective: Build a neural network to classify whether a tumor is malignant or benign using the Breast Cancer dataset.

The screenshot shows a Jupyter Notebook interface with several code cells and their execution results. The code is written in Python and follows these steps:

- Imports necessary libraries: numpy, pandas, matplotlib.pyplot, sklearn.datasets, sklearn.model_selection, and sklearn.preprocessing.
- Loads the breast cancer dataset and performs a train-test split with test_size=0.2 and random_state=42.
- Scales the features using StandardScaler.
- Defines sigmoid and sigmoid_derivative functions.
- Initializes neural network parameters (W1, b1, W2, b2) with random values scaled by 0.01.
- Defines learning rate lr = 0.01 and initializes epochs and losses lists.
- Performs forward pass: z1 = X_train @ W1 + b1, a1 = sigmoid(z1), z2 = a1 @ W2 + b2, y_pred = sigmoid(z2).
- Computes loss using Binary Cross-Entropy: loss = -np.mean(y_train * np.log(y_pred + 1e-8) + (1 - y_train) * np.log(1 - y_pred + 1e-8)).
- Performs backpropagation: dloss = y_pred - y_train, dW2 = a1.T @ dloss, db2 = np.sum(dloss, axis=0, keepdims=True), da1 = dloss @ W2.T, dz1 = da1 * sigmoid_derivative(z1), dW1 = X_train.T @ dz1.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

[181] ✓ 0.0s Python

data = load_breast_cancer()
X = data.data
y = data.target.reshape(-1,1) # 0 = malignant, 1 = benign

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Feature scaling
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

[182] ✓ 0.0s Python

def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(x):
    s = sigmoid(x)
    return s * (1 - s)

[183] ✓ 0.0s Python

input_neurons = X_train.shape[1]
hidden_neurons = 16
output_neurons = 1

np.random.seed(0)
W1 = np.random.randn(input_neurons, hidden_neurons) * 0.01
b1 = np.zeros((1, hidden_neurons))
W2 = np.random.randn(hidden_neurons, output_neurons) * 0.01
b2 = np.zeros((1, output_neurons))

[184] ✓ 0.0s Python

lr = 0.01
epochs = 5000
losses = []

for epoch in range(epochs):
    # Forward pass
    z1 = X_train @ W1 + b1
    a1 = sigmoid(z1)
    z2 = a1 @ W2 + b2
    y_pred = sigmoid(z2)

    # Loss (Binary Cross-Entropy)
    loss = -np.mean(y_train * np.log(y_pred + 1e-8) + (1 - y_train) * np.log(1 - y_pred + 1e-8))
    losses.append(loss)

    # Backpropagation
    dloss = y_pred - y_train
    dW2 = a1.T @ dloss
    db2 = np.sum(dloss, axis=0, keepdims=True)

    da1 = dloss @ W2.T
    dz1 = da1 * sigmoid_derivative(z1)
    dW1 = X_train.T @ dz1
```

```

# Loss (Binary Cross-Entropy)
loss = -np.mean(y_train*np.log(y_pred+1e-8) + (1-y_train)*np.log(1-y_pred+1e-8))
losses.append(loss)

# Backpropagation
dloss = y_pred - y_train
dW2 = a1.T @ dloss
db2 = np.sum(dloss, axis=0, keepdims=True)

da1 = dloss @ W2.T
dz1 = da1 * sigmoid_derivative(z1)
dW1 = X_train.T @ dz1
db1 = np.sum(dz1, axis=0, keepdims=True)

# Update weights
W2 -= lr * dW2
b2 -= lr * db2
W1 -= lr * dW1
b1 -= lr * db1

if (epoch+1) % 500 == 0:
    print(f"Epoch {epoch+1}, Loss: {loss:.4f}")

[5] ✓ 4.3s
Epoch 500, Loss: 0.0229
Epoch 1000, Loss: 0.0062
Epoch 1500, Loss: 0.0023
Epoch 2000, Loss: 0.0012
Epoch 2500, Loss: 0.0008
Epoch 3000, Loss: 0.0006
Epoch 3500, Loss: 0.0005
Epoch 4000, Loss: 0.0004
Epoch 4500, Loss: 0.0003
Epoch 5000, Loss: 0.0003

```

Code + Markdown | ▶ Run All ⚡ Restart 🗑 Clear All Outputs | Jupyter Variables | Outline ... Python 3.13

```

[186] Epoch 5000, Loss: 0.0003
Epoch 3500, Loss: 0.0005
Epoch 4000, Loss: 0.0004
Epoch 4500, Loss: 0.0003
Epoch 5000, Loss: 0.0003

▶ 
Success
[186] ✓ 0.0s Python
... Training Accuracy: 1.0

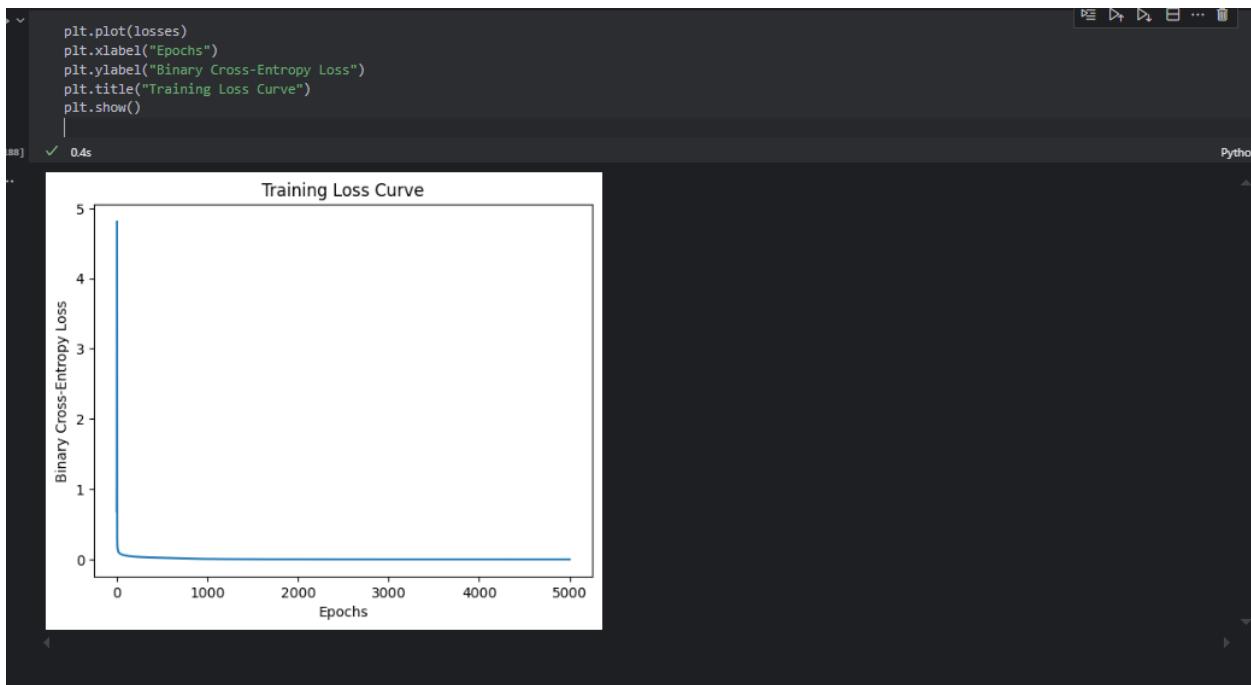
[187]
# Forward pass on test set
a1_test = sigmoid(X_test @ W1 + b1)
y_test_pred = sigmoid(a1_test @ W2 + b2)
y_test_pred_binary = (y_test_pred > 0.5).astype(int)

test_accuracy = np.mean(y_test_pred_binary == y_test)
print("Test Accuracy:", test_accuracy)

[187] ✓ 0.0s Python
... Test Accuracy: 0.9736842105263158

▶ 
plt.plot(losses)
plt.xlabel("Epochs")
plt.ylabel("Binary Cross-Entropy Loss")
plt.title("Training Loss Curve")
plt.show()

```



Question 5

Multi-Class Classification on Iris Dataset

Objective: Train a neural network to classify flower species (Setosa, Versicolor, Virginica).

```

lab5.ipynb > plt.plot(losses)
Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline ... Python 3.13.7

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
import matplotlib.pyplot as plt

# CSV file path
data = pd.read_csv("iris_data.csv") # Replace with your CSV file name

# Check first few rows
print(data.head())

```

[235] ✓ 0.0s Python

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

```

...      species
0       setosa
1       setosa
2       setosa
3       setosa
4       setosa

```

[236] ✓ 0.0s Python

```

# Features (all columns except last)
X = data.iloc[:, :-1].values

```

lab5.ipynb > plt.plot(losses)

+ Code + Markdown | ▶ Run All ⚡ Restart ⌂ Clear All Outputs | Jupyter Variables ⌂ Outline ... Python 3.13.7

[237] ✓ 0.0s

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

[238] ✓ 0.0s

```
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(x):
    s = sigmoid(x)
    return s * (1 - s)

def softmax(x):
    exp_x = np.exp(x - np.max(x, axis=1, keepdims=True))
    return exp_x / np.sum(exp_x, axis=1, keepdims=True)
```

[239] ✓ 0.0s

```
input_neurons = X_train.shape[1] # 4 features
hidden_neurons = 8
output_neurons = y_train.shape[1] # 3 classes

np.random.seed(0)
W1 = np.random.randn(input_neurons, hidden_neurons) * 0.01
b1 = np.zeros((1, hidden_neurons))
W2 = np.random.randn(hidden_neurons, output_neurons) * 0.01
b2 = np.zeros((1, output_neurons))
```

[240]

Code + Markdown | ▶ Run All ⚡ Restart ⌂ Clear All Outputs | Jupyter Variables ⌂ Outline ... Python 3.13.7

```
# Backpropagation
dloss = (y_pred - y_train)/y_train.shape[0]
dw2 = a1.T @ dloss
db2 = np.sum(dloss, axis=0, keepdims=True)

da1 = dloss @ W2.T
dz1 = da1 * sigmoid_derivative(z1)
dw1 = X_train.T @ dz1
db1 = np.sum(dz1, axis=0, keepdims=True)

# Update weights
W2 -= lr * dw2
b2 -= lr * db2
W1 -= lr * dw1
b1 -= lr * db1

if (epoch+1) % 500 == 0:
    print(f"Epoch {epoch+1}, Loss: {loss:.4f}")
```

[41] ✓ 2.7s

```
Epoch 500, Loss: 0.6236
Epoch 1000, Loss: 0.3484
Epoch 1500, Loss: 0.2495
Epoch 2000, Loss: 0.1868
Epoch 2500, Loss: 0.1478
Epoch 3000, Loss: 0.1240
Epoch 3500, Loss: 0.1086
Epoch 4000, Loss: 0.0980
Epoch 4500, Loss: 0.0903
Epoch 5000, Loss: 0.0845
```

```

# Training accuracy
train_pred = np.argmax(y_pred, axis=1)
y_train_labels = np.argmax(y_train, axis=1)
train_accuracy = np.mean(train_pred == y_train_labels)
print("Training Accuracy:", train_accuracy)

# Test accuracy
a1_test = sigmoid(X_test @ W1 + b1)
y_test_pred = softmax(a1_test @ W2 + b2)
test_pred_labels = np.argmax(y_test_pred, axis=1)
y_test_labels = np.argmax(y_test, axis=1)
test_accuracy = np.mean(test_pred_labels == y_test_labels)
print("Test Accuracy:", test_accuracy)

[242] ✓ 0.0s Python
... Training Accuracy: 0.9583333333333334
Test Accuracy: 1.0

▷ plt.plot(losses)
plt.xlabel("Epochs")
plt.ylabel("Categorical Cross-Entropy Loss")
plt.title("Training Loss Curve - Iris Dataset")
plt.show()
|
[243] ✓ 0.3s Python

```



```

[243] ✓ 0.3s Python
... plt.plot(losses)
plt.xlabel("Epochs")
plt.ylabel("Categorical Cross-Entropy Loss")
plt.title("Training Loss Curve - Iris Dataset")
plt.show()

[243] ✓ 0.3s Python

```

Question 6

Regression Problem (House Price Prediction)

Objective: Predict house prices using the **California Housing dataset**.

```
+ Code + Markdown | ⏪ Run All ⏴ Restart ⏴ Clear All Outputs | 📁 Jupyter Variables | 📄 Outline ... Python 3.13
D losses = []
for e in range(epochs):
    # Forward
    a1 = sigmoid(X_train @ W1 + b1)
    y_pred = a1 @ W2 + b2

    # Loss
    loss = np.mean((y_train - y_pred)**2)
    losses.append(loss)

    # Backprop & update
    dw2 = a1.T @ (2*(y_pred - y_train)/y_train.shape[0])
    db2 = np.sum(2*(y_pred - y_train)/y_train.shape[0], axis=0, keepdims=True)
    dz1 = (2*(y_pred - y_train)/y_train.shape[0]) @ W2.T * a1*(1-a1)
    dw1 = X_train.T @ dz1
    db1 = np.sum(dz1, axis=0, keepdims=True)

    W2 -= lr*dw2; b2 -= lr*db2
    W1 -= lr*dw1; b1 -= lr*db1

    if (e+1) % 200 == 0:
        print(f"Epoch {e+1}, Loss: {loss:.4f}")

[283] ✓ 35.8s
```

... Epoch 200, Loss: 0.9993
Epoch 400, Loss: 0.9968
Epoch 600, Loss: 0.9872
Epoch 800, Loss: 0.9511
Epoch 1000, Loss: 0.8464
Epoch 1200, Loss: 0.6709
Epoch 1400, Loss: 0.5312
Epoch 1600, Loss: 0.4710
Epoch 1800, Loss: 0.4483
Epoch 2000, Loss: 0.4356

```
Epoch 800, Loss: 0.9511
Epoch 1000, Loss: 0.8464
Epoch 1200, Loss: 0.6709
Epoch 1400, Loss: 0.5312
Epoch 1600, Loss: 0.4710
Epoch 1800, Loss: 0.4483
Epoch 2000, Loss: 0.4356

a1_test = sigmoid(X_test @ W1 + b1)
y_test_pred = a1_test @ W2 + b2

y_test_pred_orig = scaler_y.inverse_transform(y_test_pred)
y_test_orig = scaler_y.inverse_transform(y_test)

print("R² Score:", r2_score(y_test_orig, y_test_pred_orig))
print("MSE:", mean_squared_error(y_test_orig, y_test_pred_orig))

[283] ✓ 0.0s
```

... R² Score: 0.545118164526881
MSE: 0.5960812412426748

```
D plt.scatter(y_test_orig, y_test_pred_orig, alpha=0.5)
plt.plot([y_test_orig.min(), y_test_orig.max()], [y_test_orig.min(), y_test_orig.max()], 'r--')
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual vs Predicted House Prices")
plt.show()

[284] ✓ 0.7s
```



Question 7

Neural Network with Dropout Regularization

Objective: Prevent overfitting using Dropout layers on the MNIST digit dataset.

```

Code + Markdown | Run All ⚡ Restart ⚡ Clear All Outputs | Jupyter Variables | Outline ...
Python 3.13.7

98] ✓ 0.0s

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Flatten
from tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

# Load MNIST
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# Normalize images
x_train, x_test = x_train/255.0, x_test/255.0

# One-hot encode labels
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

99] ✓ 1.3s

model = Sequential([
    Flatten(input_shape=(28,28)),      # Flatten 28x28 images
    Dense(128, activation='relu'),
    Dropout(0.3),                   # Dropout 30%
    Dense(64, activation='relu'),
    Dropout(0.3),                   # Dropout 30%
    Dense(10, activation='softmax')   # Output layer for 10 classes
])
100] ✓ 0.1s

```

```
- Code + Markdown | Run All ⏪ Restart ⏴ Clear All Outputs | Jupyter Variables ⏴ Outline ... Python 3.13.7

model = Sequential([
    Flatten(input_shape=(28,28)),      # Flatten 28x28 images
    Dense(128, activation='relu'),    # Dropout 30%
    Dropout(0.3),
    Dense(64, activation='relu'),
    Dropout(0.3),
    Dense(10, activation='softmax')   # Output layer for 10 classes
])

[380] ✓ 0.1s Python

model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

[381] ✓ 0.0s Python

> history = model.fit(
    x_train, y_train,
    validation_split=0.2,
    epochs=20,
    batch_size=128,
    verbose=1
)

[382] ✓ 1m 22.0s Python

... Epoch 1/20
375/375 7s 10ms/step - accuracy: 0.8199 - loss: 0.5910 - val_accuracy: 0.9426 - val_loss: 0.1983
Epoch 2/20
375/375 3s 8ms/step - accuracy: 0.9249 - loss: 0.2614 - val_accuracy: 0.9582 - val_loss: 0.1476
Epoch 3/20
```

```
lab1.pynd > history = model.fit(
+ Code + Markdown | Run All ⏪ Restart ⏴ Clear All Outputs | Jupyter Variables ⏴ Outline ... Python 3

    loss, accuracy = model.evaluate(x_test, y_test)
    print("Test Loss:", loss)
    print("Test Accuracy:", accuracy)

[383] ✓ 1.7s Python

... 313/313 2s 5ms/step - accuracy: 0.9799 - loss: 0.0717
Test Loss: 0.071674905173729
Test Accuracy: 0.9799000024795532
```