

Lab No. 13

Natural Language Processing (NLP)

This laboratory session introduces students to Word2Vec, a popular technique in Natural Language Processing (NLP) used to represent words as meaningful dense vectors. Using textual data from the *Game of Thrones* series, students will learn how word embeddings capture semantic relationships between words based on their context. Through preprocessing, model training, and similarity analysis, this lab helps students understand how machines learn word meanings and relationships from large text corpora in an unsupervised manner.

LAB Objectives:

- Understand **distributional semantics**
- Learn how **Word2Vec** converts words into dense vectors
- Apply **Word2Vec (Skip-Gram / CBOW)** on real textual data (*Game of Thrones*)
- Explore **word similarity, analogy, and visualization**

game-of-thrones-word2vec

word2vec applied on game of thrones data

Dataset Link: <https://www.kaggle.com/khulasasndh/game-of-thrones-books>

Download the data set from Kaggle

Game Of Thrones books

Data Card Code (47) Discussion (0) Suggestions (0)

29 Code Download

001ssb.txt (1.63 MB)

Download

About this file

This file does not have a description yet.

Suggest Edits

 This preview is truncated due to the large file size. Create a Notebook or download this file to see the full content.

Download

Create Notebook

A Game Of Thrones

Book One of A Song of Ice and Fire

By George R. R. Martin

PROLOGUE

"We should start back," Gared urged as the woods began to grow dark around them. "The wildlings are dead."

Data Explorer

Version 1 (9.9 MB)

- 001ssb.txt
- 002ssb.txt
- 003ssb.txt
- 004ssb.txt
- 005ssb.txt

Summary

5 files

Add the dataset in folder data where VS code directory present

VScode Examples				
File	Home	Share	View	
This PC > Local Disk (C:) > Users > Syed Hamedoon > VScode Examples				
Name	Date modified	Type	Size	
.vscode	8/25/2025 3:54 PM	File folder		
data	12/22/2025 12:03 PM	File folder		
myenv	9/10/2025 10:10 AM	File folder		
archive (1)	12/22/2025 12:02 PM	WinRAR ZIP archive	3,801 KB	
ArraysP	9/10/2025 10:33 AM	Python Source File	1 KB	
cardpolarrunner	12/17/2025 11:55 AM	Jupyter Source File	1 KB	
cartpolarun	12/16/2025 4:33 PM	Python Source File	1 KB	
chatbot	12/14/2025 2:11 PM	Python Source File	1 KB	
Churn_Modelling	10/27/2025 9:51 AM	Microsoft Excel Co...	669 KB	
DataFile	9/24/2025 12:21 PM	Microsoft Excel Co...	2 KB	
DTandReg	10/8/2025 9:35 AM	Jupyter Source File	76 KB	
DThandwriting	10/8/2025 10:11 AM	Jupyter Source File	57 KB	
FeaturePins	10/20/2025 8:51 AM	Python Source File	8 KB	
inputs	9/9/2025 2:40 PM	Python Source File	1 KB	
Iris	9/21/2019 5:26 PM	Microsoft Excel Co...	5 KB	

Code in jupiter VS code:

```
import numpy as np  
import pandas as pd
```

```
!pip install gensim
```

```
import gensim  
import os
```

```
!pip install nltk
```

```
data = "C:/Users/Syed Hamedoon/VScode Examples/data"
```

```
import nltk  
  
nltk.download('punkt')  
nltk.download('punkt_tab')
```

```
import os  
from nltk import sent_tokenize  
from gensim.utils import simple_preprocess  
  
DATA_PATH = r"C:\Users\Syed Hamedoon\VScode Examples\data"  
  
story = []  
  
for filename in os.listdir(DATA_PATH):  
    if filename.endswith(".txt"):  
        file_path = os.path.join(DATA_PATH, filename)  
  
        try:  
            with open(file_path, "r", encoding="utf-8") as f:  
                corpus = f.read()  
        except UnicodeDecodeError:  
            with open(file_path, "r", encoding="cp1252") as f:  
                corpus = f.read()
```

```
for sent in sent_tokenize(corpus):
    story.append(simple_preprocess(sent))
```

```
print(len(story))
print(story[:2])
```

```
model = gensim.models.Word2Vec(
    window=10,
    min_count=2
)
```

```
model.build_vocab(story)
```

```
model.train(story, total_examples=model.corpus_count, epochs=model.epochs)
```

```
model.wv.most_similar('daenerys')
```

```
model.wv.doesnt_match(['jon','rikon','robb','arya','sansa','bran'])
```

```
model.wv.doesnt_match(['cersei', 'jaime', 'bronn', 'tyrion'])
```

```
model.wv['king']
```

```
model.wv.similarity('arya','sansa')
```

```
model.wv.similarity('tywin','sansa')
```

```
model.wv.get_normed_vectors()
```

```
y = model.wv.index_to_key
```

```
len(y)
```

```
y
```

```
from sklearn.decomposition import PCA
```

```
pca = PCA(n_components=3)
```

```
X = pca.fit_transform(model.wv.get_normed_vectors())
```

```
X.shape
```

```
!pip install --upgrade nbformat
```

```
import pandas as pd
import plotly.express as px
import plotly.io as pio

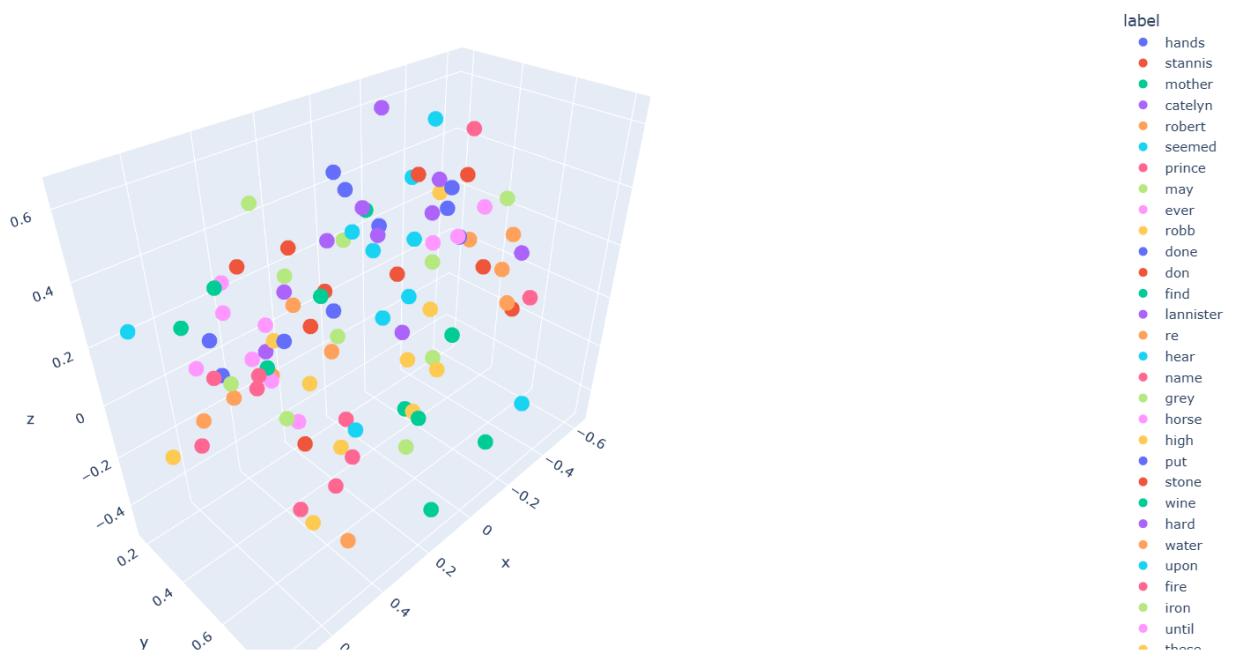
pio.renderers.default = "browser" # ← IMPORTANT

df = pd.DataFrame(X[200:300], columns=["x", "y", "z"])
df["label"] = y[200:300]

fig = px.scatter_3d(
    df,
    x="x",
    y="y",
    z="z",
    color="label"
)

fig.show()
```

Output:



Code:

```
import numpy as np
import pandas as pd
import gensim
import os
import nltk

[1] ✓ 8.0s

▷ ▾
nltk.download('punkt')
nltk.download('punkt_tab')

[2] ✓ 11.9s

... [nltk_data] Downloading package punkt to
[nltk_data]   C:\Users\h\AppData\Roaming\nltk_data...
[nltk_data]   Unzipping tokenizers\punkt.zip.
[nltk_data] Downloading package punkt_tab to
[nltk_data]   C:\Users\h\AppData\Roaming\nltk_data...
[nltk_data]   Unzipping tokenizers\punkt_tab.zip.

... True

DATA_PATH = r"C:\Users\h\Desktop\AI\data"

[3] ✓ 0.0s

from nltk import sent_tokenize
from gensim.utils import simple_preprocess

story = []

[4] ✓ 13.6s

story = []

for filename in os.listdir(DATA_PATH):
    if filename.endswith(".txt"):
        file_path = os.path.join(DATA_PATH, filename)

        try:
            with open(file_path, "r", encoding="utf-8") as f:
                UnicodeDecodeError
        except UnicodeDecodeError:
            with open(file_path, "r", encoding="cp1252") as f:
                corpus = f.read()

        for sent in sent_tokenize(corpus):
            story.append(simple_preprocess(sent))

    print(len(story))
    print(story[:2])

[5] ✓ 0.0s
145020
[['game', 'of', 'thrones', 'book', 'one', 'of', 'song', 'of', 'ice', 'and', 'fire', 'by', 'george', 'martin', 'prologue', 'we', 'should', 'start', 'back', 'gared', 'urged', 'as', 'the', 'woods', 'be

model = gensim.models.Word2Vec(
    window=10,
    min_count=2
)

[6] ✓ 0.0s
```

```

    model.build_vocab(story)

[✓] 2.9s

model.train(
    story,
    total_examples=model.corpus_count,
    epochs=model.epochs
)

[✓] 13.0s
(6569050, 8628190)

model.wv.most_similar('daenerys')

[✓] 0.0s
[('stormborn', 0.7747085690498352),
 ('myrcella', 0.7417674660682678),
 ('princess', 0.7254637479782104),
 ('targaryen', 0.7166107296943665),
 ('viserys', 0.6705787777900696),
 ('unburnt', 0.6640740633010864),
 ('margaery', 0.6635355949401855),
 ('queen', 0.6590381264686584),
 ('dorne', 0.6493701934814453),
 ('prince', 0.6415064930915833)]

model.wv.doesnt_match(['jon','rikon','robb','arya','sansa','bran'])

[✓] 0.0s

[✓] 0.0s
'jon'

model.wv.doesnt_match(['cersei', 'jaime', 'bronn', 'tyrion'])

[✓] 0.0s
'bronn'

model.wv['king']

[✓] 0.0s
array([ 0.26891482,  0.6001591 ,  2.6924682 ,  1.9977348 , -2.0928023 ,
       1.6741775 , -0.74811924,  1.1452627 , -1.9751002 ,  0.2568304 ,
      -1.0746068 ,  0.86136997,  0.500974 ,  2.8437057 , -3.4520917 ,
      -1.4782377 ,  1.6148216 ,  2.1974611 ,  0.72354144,  1.636202 ,
       0.8321633 , -0.9486615 ,  2.9643846 , -3.2509696 , -2.0574193 ,
      2.2818778 , -0.37232143, -1.0752884 , -0.47045296,  1.9520614 ,
     -2.7735584 , -0.2271765 ,  0.5828542 , -0.68766874,  1.1737362 ,
     -1.5275186 , -0.4710071 , -1.7312647 ,  0.64836687, -2.2143753 ,
       0.1098459 ,  1.5535835 ,  1.6918998 ,  0.23302521, -0.33263117,
     -1.9727484 ,  0.02945553, -1.9833945 ,  2.4889457 , -3.0467474 ,
     -2.7870016 , -2.498347 , -1.981903 , -2.5230064 ,  3.9048395 ,
     -2.4079874 ,  1.5737748 ,  1.081367 ,  1.0208126 ,  1.1831667 ,
      1.2238128 ,  1.3786469 ,  0.9803622 ,  1.2135874 , -0.08005533,
      1.9667065 , -0.71088463, -1.6886363 , -1.1306278 , -1.1948683 ,
     -0.2185671 , -0.37058043,  3.992105 , -2.2785435 ,  2.4370081 ,
     -0.26333271 , -1.77131969, -1.12305627 , -1.8899138 , -3.097366 ,
      1.600244 ,  2.0607677 , -1.9047579 , -2.1947439 , -2.5791237 ,
     -3.3174608 ,  0.8148284 ,  4.4358478 ,  1.1345456 , -0.6091975 ,
      1.0613319 ,  1.52664478 , -0.29228923, -1.6843776 ,  0.23084323,
     -3.454302 , -1.7525158 , -0.992976 , -1.433621 ,  0.1680581 ],
      dtype=float32)

```

```
model.wv.similarity('arya', 'sansa')

6] ✓ 0.0s
np.float32(0.8439563)

model.wv.similarity('tywin', 'sansa')

7] ✓ 0.0s
np.float32(0.254941)

vectors = model.wv.get_normed_vectors()
words = model.wv.index_to_key

len(words)

8] ✓ 0.0s
17453

from sklearn.decomposition import PCA

pca = PCA(n_components=3)
X = pca.fit_transform(vectors)

X.shape

9] ✓ 1.0s
(17453, 3)

from sklearn.decomposition import PCA

pca = PCA(n_components=3)
X = pca.fit_transform(vectors)

X.shape

10] ✓ 1.0s
(17453, 3)

import plotly.express as px
import plotly.io as pio

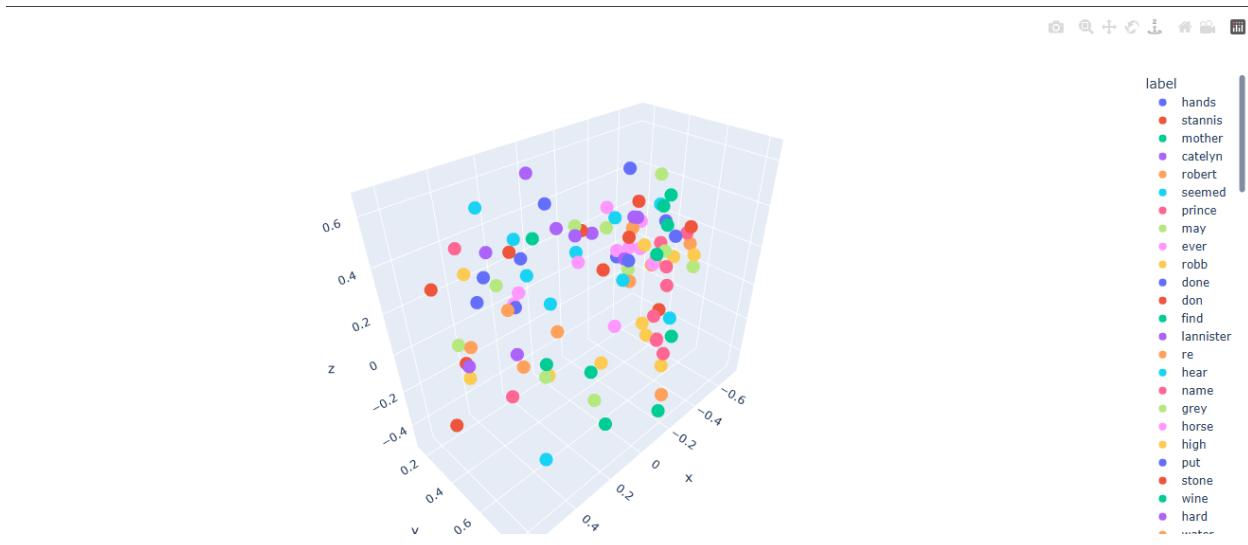
pio.renderers.default = "browser"

df = pd.DataFrame(X[200:300], columns=["x", "y", "z"])
df["label"] = words[200:300]

fig = px.scatter_3d(
    df,
    x="x",
    y="y",
    z="z",
    color="label"
)
fig.show()

11] ✓ 9.1s
```

Output:



LAB Questions

Q1.

What is the core idea behind Word2Vec?

Answer:

Word2Vec is a technique that converts words into dense vectors which capture both semantic and syntactic relationships. The main idea is that words appearing in similar contexts have similar meanings. It learns vector representations such that similar words are close together in the vector space.

Q2.

Difference between CBOW and Skip-Gram?

Answer:

Feature	CBOW	Skip-Gram
Goal	Predict target word from context	Predict context words from target
Input	Context words	Target word
Output	Target word	Context words
Speed	Faster, good for frequent words	Slower, better for rare words
Use Case	Frequent words	Rare words

Q3.

Why is one-hot encoding inefficient?

Answer:

One-hot encoding creates very large and sparse vectors that do not show any similarity between words. Every word is independent in this representation, and it requires a lot of memory and computation for large vocabularies.

Q4.

Why do character names appear close in vector space?

Answer:

- Character names often appear in similar contexts like same scenes or dialogues.
- Word2Vec places them close together because it learns similarity from context.

Q5.

How does window size affect semantic learning?

Answer:

Window size determines how many words around the target word are considered while training. A small window focuses more on syntactic relations like grammar, while a large window captures broader semantic relations like topics. If the window is too large, irrelevant words may add noise, while a very small window may miss context.

Q6.

Why might rare characters have poor embeddings?

Answer:

Rare words or characters appear infrequently in the text. Because Word2Vec relies on word co-occurrence in contexts to learn embeddings, rare words get very little information, leading to vectors that do not accurately capture their meaning.

Q7.

Which model performed better: CBOW or Skip-Gram? Why?

Answer:

Skip-Gram usually performs better than CBOW, especially for rare words. It predicts multiple context words for each target word, so even rare words get more training examples. CBOW is faster and works well for frequent words, but it may not handle rare words effectively.

Q8.

What happens if vector size is too small or too large?

Answer:

- Too small vector size cannot capture enough information, making embeddings less useful.
- Too large vector size may overfit the training data, capture noise, and require more computation.
- Choosing an optimal size, usually between 100 to 300 dimensions, balances expressiveness and efficiency.

Q9.

Can Word2Vec understand word meaning without labels? Explain.

Answer:

Yes, Word2Vec works in an unsupervised way. It learns the meaning of words purely from the contexts they appear in, without any predefined labels. Semantic similarity is derived from co-occurrence patterns in the text.